

The Long Tail of Indie SaaS

What 1,000 TrustMRR listings reveal about a marketplace that looks bigger than it is



Every company in the dataset, ranked by revenue — the shape of the whole argument

What this deck is actually testing

“TrustMRR aggregates 1,000 self-reported indie / micro-SaaS startups with public revenue dashboards. Does this dataset describe a healthy, diverse startup ecosystem — or a highly concentrated, US-and-Stripe-dependent market where a small number of outliers explain most of the revenue and most listed companies are stagnant or dying? This presentation uses the dataset itself to test that question.”

A

A healthy, diverse ecosystem

Revenue spread across many companies, many countries and many payment rails. Growth is broad-based. The average listing is a reasonable benchmark.

B

A concentrated long-tail market

A few breakouts hold most of the revenue. The typical listing is tiny and flat, clustered in one country and on one payment provider.

The next seven slides test A against B, one dimension at a time.

The typical company is tiny. The average one doesn't exist.

30-DAY REVENUE

\$400

median | mean \$15,596 — 39× apart

ALL-TIME REVENUE


\$3,620

median | mean \$334,044 — 92× apart

CONCENTRATION

49.6%


of all 30-day revenue comes from 10 of the 1,000 listings



- The mean is not a usable “typical company” number here — a handful of outliers (largest: \$76.6M all-time) drag it 39× above the median. Report both, always, and lead with the median.

All 1,000 rows. Median and mean computed over the full dataset; the log-scale histogram omits the 216 listings reporting \$0 in the last 30 days, which a log axis cannot show.

Growth is mostly an illusion of the average



MEAN 30-DAY MRR GROWTH

+33.1%

Sounds like a healthy, expanding market.

MEDIAN 30-DAY MRR GROWTH

0.0%

The middle company in this dataset did not grow at all.

**61.9% of listings are flat or shrinking.
Only 27.4% grew.**

● **“Average growth” is being carried by a few rocketships. Strip them out and the middle of this marketplace is standing still — or going backwards.**

Field: mrr_growth_30d_pct. Shares are of all 1,000 listings; the 10.7% (107 rows) with no reported growth value are shown as their own band rather than silently dropped. “Exactly flat” means a reported value of precisely 0%.

Globally listed. American in revenue.

UNITED STATES


25.8%

of listings — but 63.1% of all 30-day revenue

NEXT LARGEST — FRANCE

\$669K

30-day revenue: under 7% of the US total, from 7.9% of listings




24.1% of rows carry no country at all — and those rows hold 17.6% of the revenue, so the geographic picture is genuinely incomplete, not just sparse.

● By headcount TrustMRR looks like a global marketplace. By revenue it is an American one with international spectators.

Country field, all 1,000 rows; unlisted rows shown as their own band rather than excluded. Two rows carry a corrupted boolean value (FALSE) where a country code should be — they hold \$1,709 combined and are excluded from named-country figures.

AI wins on volume. It does not win per company.



- **AI is the biggest by both counts**

197 listings (the largest category) and \$2.36M in 30-day revenue (the largest total).

- **...and near the bottom per company**

Median MRR of \$446 — below Fintech, Social Media, Content Creation, Analytics, Marketing and Education.

- **Small categories punch above their weight**

Education: 30 listings, \$2,142 median MRR. Marketing: 42 listings, \$1,586. Both beat AI per company by 3–5x.

- **Volume and hype are not unit economics. The categories with the strongest typical company are the small, unglamorous ones — and “AI” is where the crowd is, not where the money-per-company is.**

Categories with at least 15 listings shown (13 of 30). 27.4% of all rows have no category and are excluded from this view. Bars are medians; category totals are separately dominated by one or two outliers each.

Two-thirds of this ecosystem runs on one vendor

STRIPE


64.8%

of single-provider listings run on Stripe and nothing else

REVENUECAT

19.3%

a distant second — largely mobile-app subscription billing



Everything else — Dodo, Polar, Lemon Squeezy, Superwall, Paddle, Whop, Creem — shares the remaining 15.9%.

- **This is a Stripe-centric ecosystem. One vendor's pricing change, policy shift or outage would land on roughly two-thirds of the businesses listed here at the same time.**

Single-provider listings only (n = 981 of 1,000). The 19 listings reporting two or more providers are excluded so that shares sum cleanly; counting Stripe anywhere in the field instead gives 64.9% of all 1,000 rows — the conclusion is unchanged.

Priced like side businesses, not like startups

LISTINGS FOR SALE

26.2%

262 of 1,000 listings are actively on the market

MEDIAN ASKING PRICE

\$35,000

the midpoint of what a listed business asks for itself

MEDIAN REVENUE MULTIPLE

2.3×

what buyers are being asked to pay per unit of revenue

MEAN REVENUE MULTIPLE

7.8×

the same outlier distortion, now in the pricing layer

For scale: a \$35,000 asking price against a \$400 median 30-day revenue is roughly seven years of revenue at the median — which is why the multiple, not the price, is the number to argue about. The 2.3× median is small-business territory; the 7.8× mean is what happens when a handful of listings ask for venture-style pricing on the same page.

● **This is not a market of aspiring unicorns. It prices like what it is: a market for small, cash-flowing side businesses.**

On-sale listings only (n = 262). asking_price_usd and revenue_multiple are structurally missing for the 73.8% of listings not on sale — that missingness is by design, not a data fault. One on-sale listing has no asking price.

What the data won't tell you

ANONYMOUS LISTINGS

13.0%

130 listings trade as “Stealth Company” or “Hidden Business” — on a platform named TrustMRR

DOMAIN AUTHORITY VS REVENUE

$r = 0.278$

weak linear correlation (Spearman $\rho = 0.516$). SEO strength is not a reliable predictor of revenue here.

Everything on this platform is self-reported, with no audit trail. Treat each figure as a claim, not a measurement.

MISSINGNESS, DECLARED HONESTLY

revenue_per_visitor_usd	85.8%
visitors_30d	82.2%
asking_price_usd (structural)	73.9%
revenue_multiple (structural)	73.8%
domain_rating	28.6%
category	27.4%
country	24.1%
mrr_growth_30d_pct	10.7%

Plus: 2 rows carry a boolean FALSE in the country field — a type error, not a missing value.

- **Two of the most analytically interesting fields are more than 80% empty, and an eighth of the companies are anonymous. Any conclusion drawn here is a conclusion about what founders chose to publish.**

Correlations computed on the 714 rows carrying a domain_rating. Pearson is depressed by revenue outliers; the higher rank correlation ($\rho = 0.516$) says the ordering holds better than the linear fit — the gap between the two is itself evidence of the long tail.

THE VERDICT

A concentrated long-tail market — not a diverse ecosystem

Every dimension pointed the same way:

- Revenue: 10 of 1,000 listings hold half of it
- Growth: the median company grew 0%
- Geography: 25.8% of listings, 63.1% of revenue
- Infrastructure: 64.8% on a single payment vendor
- Pricing: a 2.3× median multiple on a \$35,000 ask

THE ONE NUMBER TO REMEMBER

\$400 vs \$15,596

Median vs mean 30-day revenue. If you only keep one figure, keep this one — it reframes every other number in the deck.



So what? If you are benchmarking yourself against TrustMRR, benchmark against the median, not the average — the average describes nobody. If you are shopping for an acquisition, the 2.3× median multiple is the real market; the long tail is where the listings are, not the revenue.