SLIDE 1 · THE DATA

One flat table of 1,000 businesses — and a visible defect rate

1,000
Businesses in the table
29 fields, one flat sheet

262
Of them listed for sale
$86.8M asked in total

80%
Of those 262 for-sale listings
trip at least one screen rule

52
For-sale listings that trip
no rule at all

The table holds seven entities

Business	identity, country, category
Financials	revenue, MRR, growth, subs
Listing	on_sale, price, multiple
Description	value prop, pricing, persona
Payment providers	many-to-many
Markets	many-to-many (JSON)
Countries	lookup

The screen: 12 rules on price, business health and data trust

Seven most-triggered rules — count of businesses, out of all 1,000

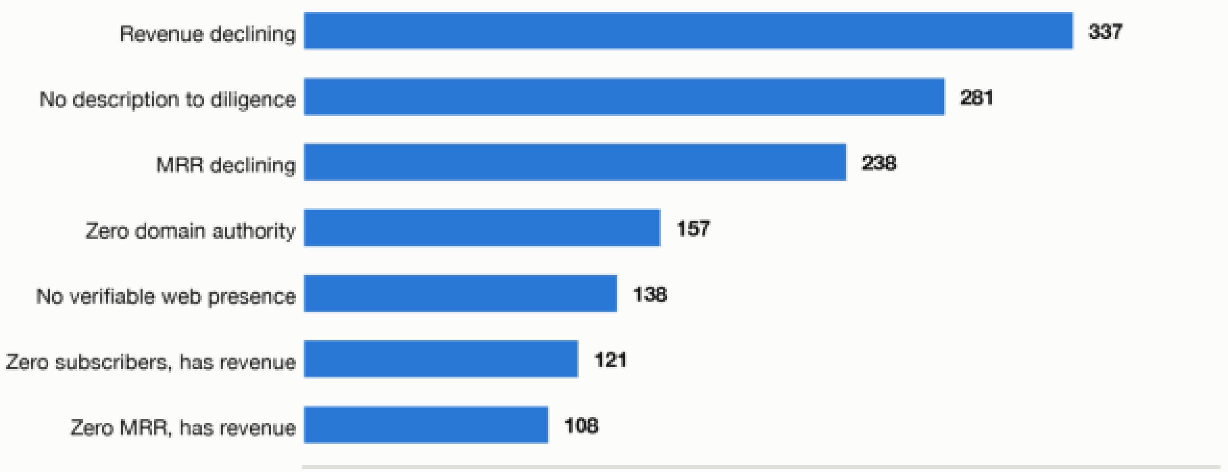

Source: TrustMRR 1,000-row dataset, frozen 30 Jul 2026. Rules read only fields already in the table. Two multi-valued columns (payment_provider, markets_json) drive the many-to-many splits.

The market does not price the defect

Clean listings ask $37k against $35k for flagged ones — 6% more — yet carry 3.3× the MRR, 3.6× the subscribers and 2.1× the domain authority.

Each panel compares median values across the 262 for-sale listings. Clean = trips no rule (52) · Flagged = trips one or more (210) · Severe = trips three or more (87)

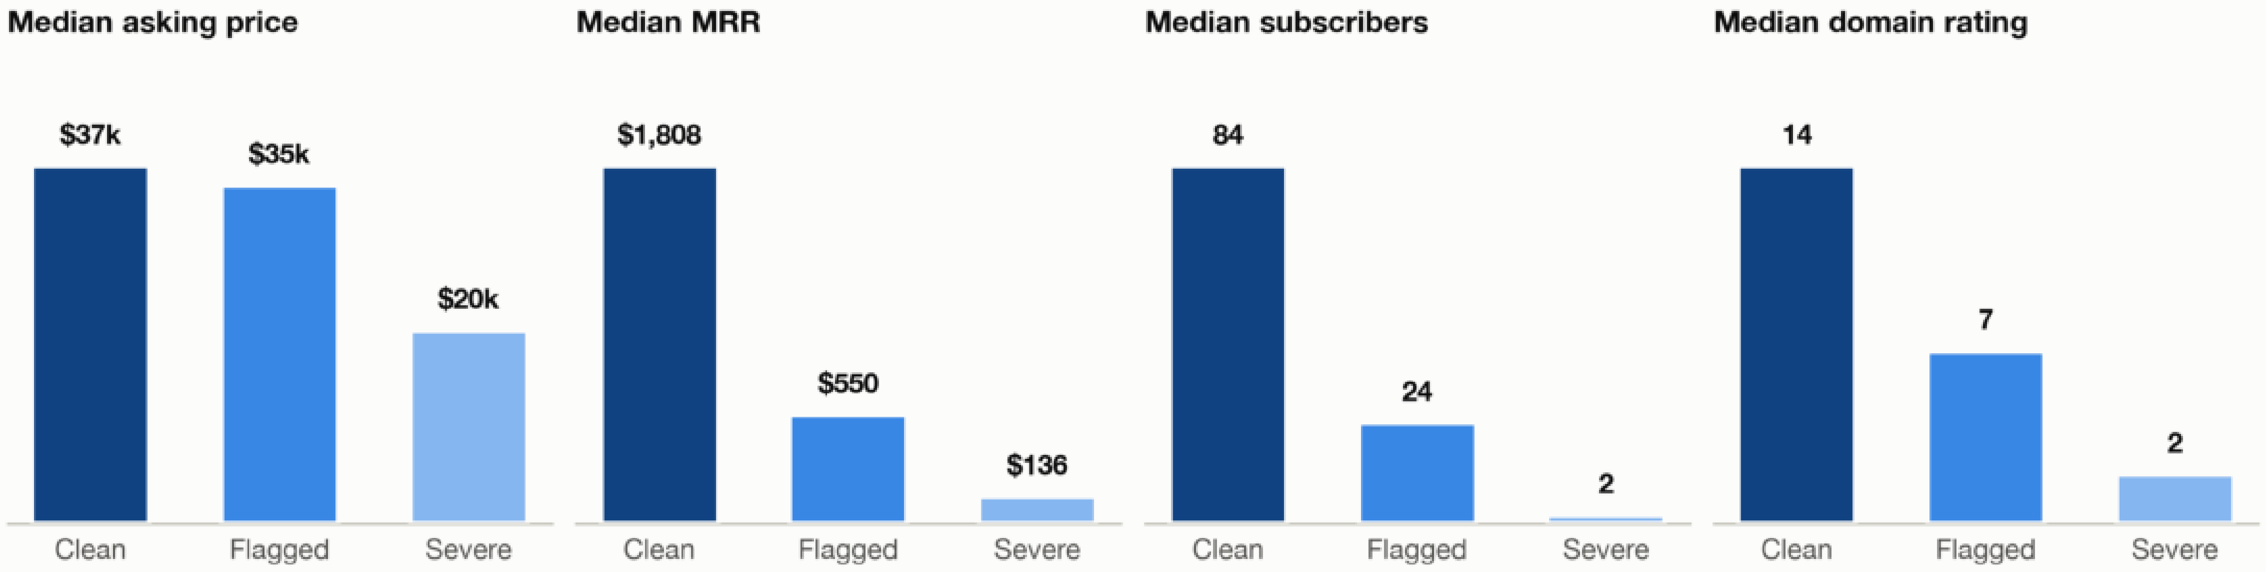

Flagged listings also carry the higher multiple — a 2.4× median against 1.75× for clean ones. A buyer pays more per dollar of revenue for the weaker asset.

Medians, not means: the pool contains a handful of $1M+ listings that would distort an average.

The thesis rests on a dataset with holes in it

How often each field is simply empty

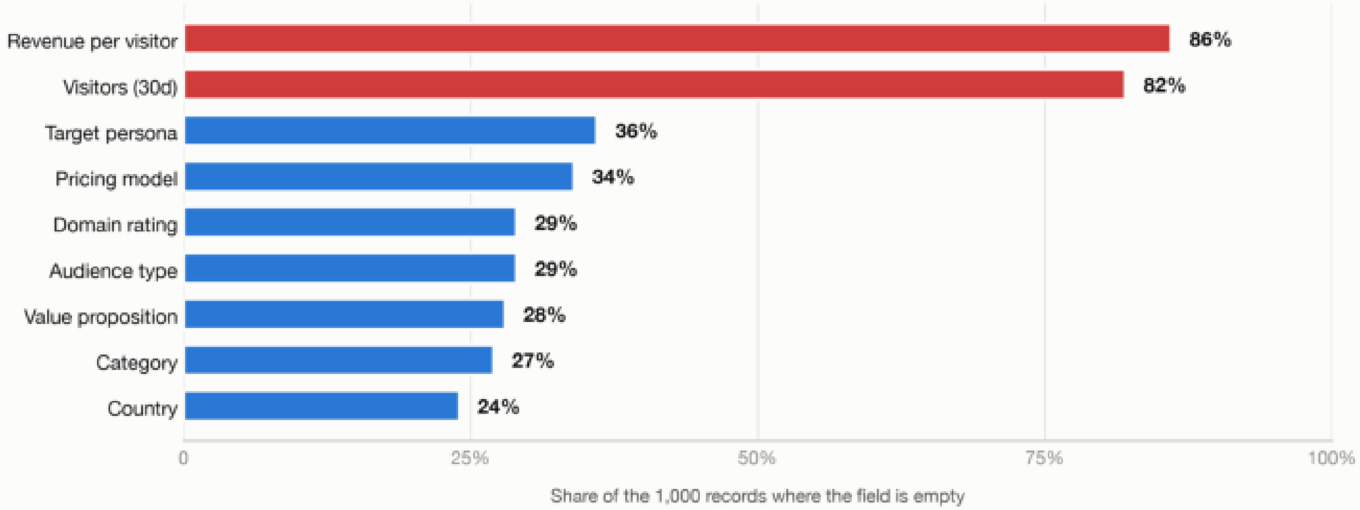

Three limits worth stating

- **Revenue cannot be cross-checked**

  Traffic is the one independent test of whether revenue is real. It is missing for 82% of records, so the screen cannot run it.

- **Flags are counted, not weighted**

  A 5× multiple currently scores the same as a missing website. Weighting would move the $28.0M figure.

- **One frozen snapshot**

  Dated 30 Jul 2026, with no history. Growth reads as a single 30-day number, so a good month and a good trend look identical.

Net effect: the 80% defect rate is a floor, not a ceiling. Fields that would catch more defects are the ones most often empty.

Missingness computed across all 1,000 records.

A clean screen still leaves a thin, small-ticket pipeline

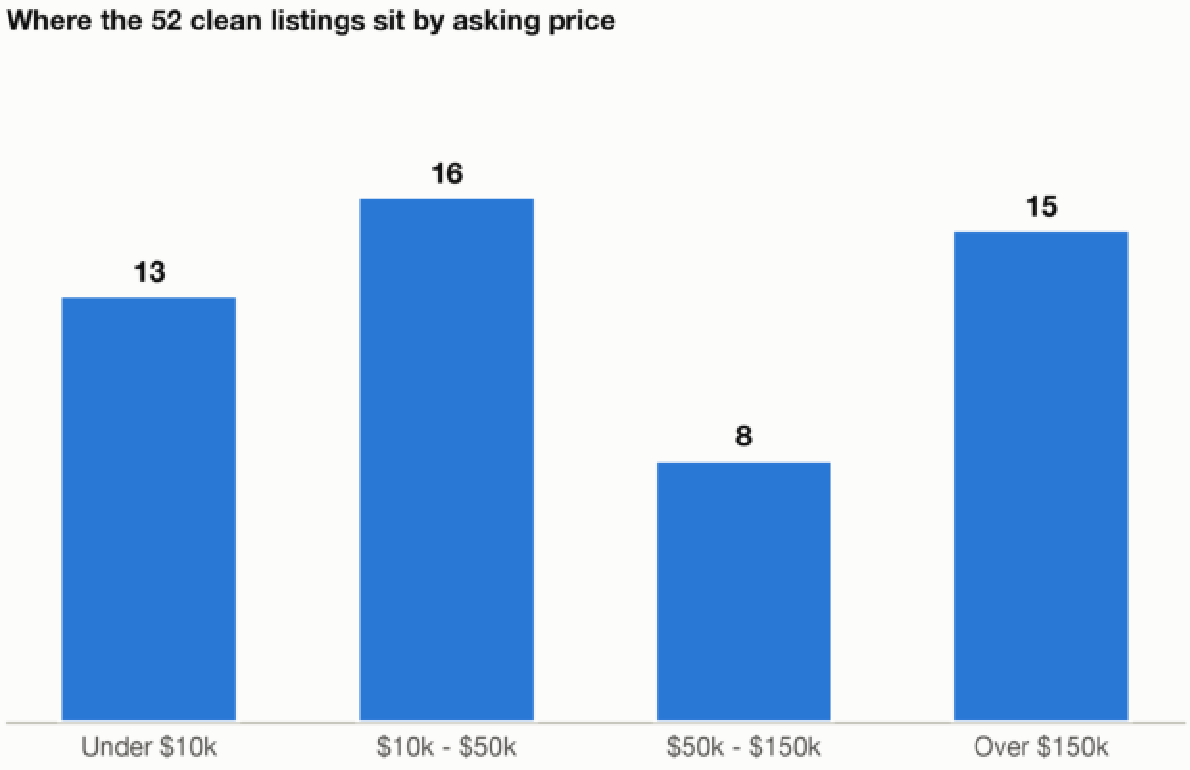

The pipeline is 52 listings

Screening 1,000 businesses yields 52 candidates. Any further filter — category, geography, ticket size — cuts into single digits fast.

More than half sit under $50k

29 of 52 clean listings are below $50k, and 13 below $10k. Diligence and transfer costs consume a large share of a deal that size.

Clean data is not a clean business

The screen tests what is reported, not what transfers. Key-person dependency, platform lock-in and undisclosed churn all survive it.

Concentration, and a blind spot

AI is the most common named category among for-sale listings (54 of 262), so a category correction hits much of the pool. A further 53 record no category at all.

52 clean listings · $9.3M combined asking value · median $37k · 29 below $50k.

SLIDE 5 · RISK AND RETURN

The screen is the edge

What the screen removes, stage by stage

$28.0M
Asking value screened out (listings tripping 3+ rules)

32%
That $28.0M as a share of the $86.8M asked

13.3×
Median MRR of clean versus severe listings

1.75×
Median revenue multiple paid on clean listings

The reward

The defect is visible in fields the platform already publishes, but it is not reflected in price. A buyer who screens pays a 6% premium and a lower multiple for roughly triple the recurring revenue, while sidestepping the $28.0M of asking value concentrated in listings that fail on three or more independent measures at once.

The edge is not picking winners. It is refusing to buy from a pool where four in five for-sale listings trip a rule.

All figures derived from the frozen dataset; see TrustMRR Risk Flags.xlsx for the per-listing screen.