STARTUP DATA DIVE · TRUSTMRR

The Fields Nobody Reads

Three signals the TrustMRR 1,000 rows gives away in columns nobody asked about

1,000

startups, payment-API verified

$133.8M

combined trailing 12-month revenue

29

columns per record

EXECUTIVE SUMMARY

The metadata is more honest than the metrics

Summary around payment rail, domain authority, and the marketplace's own pricing math

01

6.2×

The payment rail is a free classifier

Stripe subscribers pay 6.2× more per month than RevenueCat subscribers and the gap survives even for the 288 listings with no audience_type label at all

02

0.49 vs 0.17

Earned authority beats vanity reach

Domain rating predicts MRR (ρ=0.49) far better than X followers (ρ=0.17). Revenue climbs in a straight line with SEO authority but it does not climb with social following

03

24%

Multiples lie below $200 MRR

The marketplace's stated revenue multiple looks disconnected from asking price (r=0.06) until the bottom quartile of micro-revenue listings is removed, and it jumps to ρ=0.80

The average startup here describes nobody

$16,777

Mean monthly recurring revenue

$1,036

Median monthly recurring revenue

82.2%

of startups earn less than the mean

Mean sits 16.2× above median. A handful of outsized winners are doing all the work of pulling the “average” away from where most of the market actually lives

49.1% earn under $1,000 a month

26.4% clear $10,000 a month

48.0% of all MRR sits with 10 firms

Cohort: 803 startups reporting MRR above zero. Combined MRR $13.47M

Seventeen firms hold more revenue than the other 786

Cumulative share of $132.9M trailing 12-month revenue, firms ranked largest first (n=803)

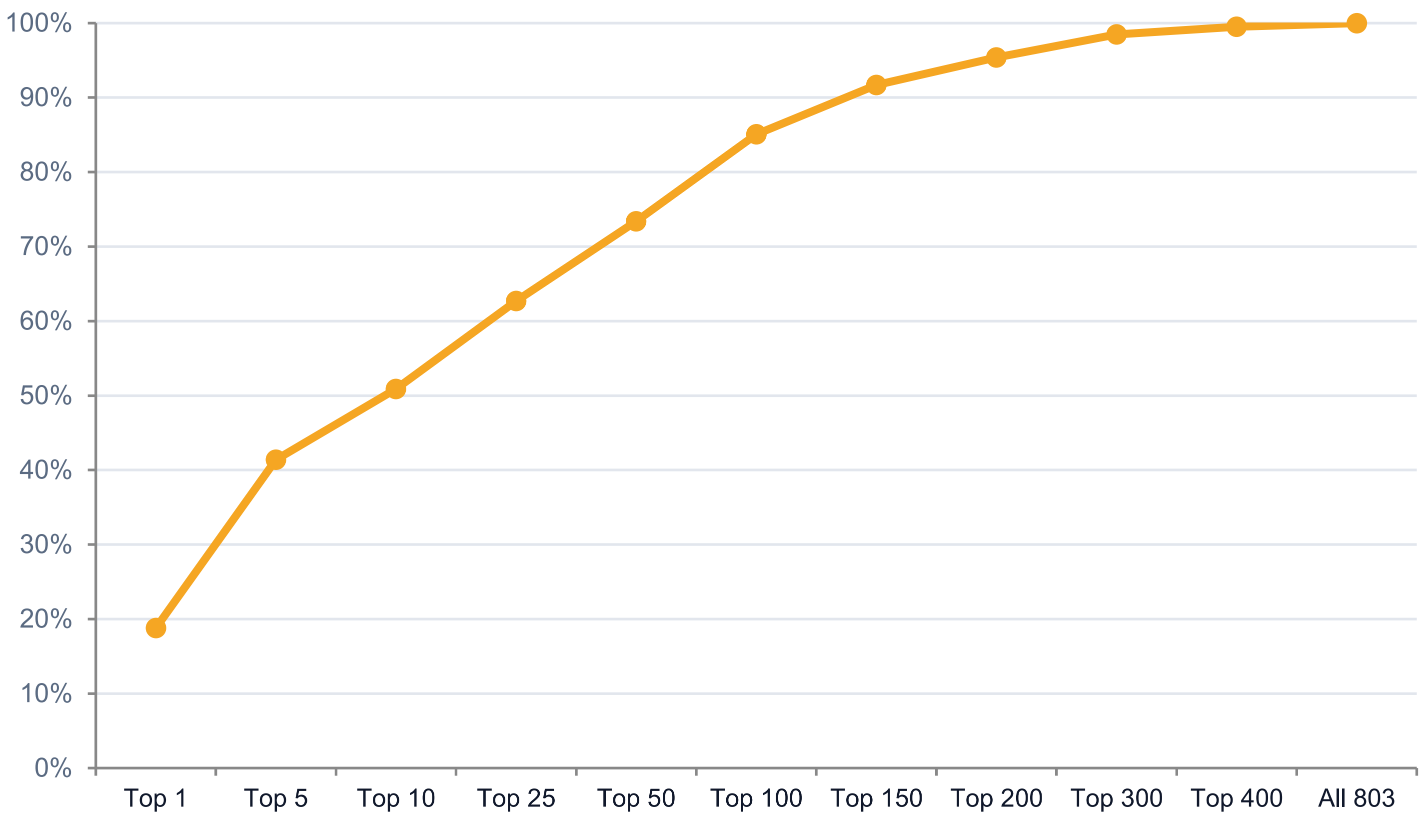

57.7%

of revenue held by the 17 firms earning above $1M a year

85.1%

held by the top 100 firms (just 12.5% of the cohort)

0.05%

held by the 222 firms earning under $1,000 a year

Cohort: 803 startups with MRR above zero. Ranked by trailing 12-month revenue

1

THE PAYMENT RAIL IS A FREE CLASSIFIER

Businesses can leave audience_type blank. They can't hide who processes their payments

audience_type is missing for 288 of 1,000 listings accounting for 28.8% of the directory **payment_provider** is missing for none of them. And the rail a founder wires up turns out to encode the same B2B/B2C economics the audience label was meant to capture completely, for every single row

$30.78

Stripe’s median $ / subscriber / mo

$5.00

RevenueCat: median $ / subscriber / mo

P.S: 677 listings with active, paying subscribers used for the per-row price comparison (Stripe n=517, RevenueCat n=160)

1 EVIDENCE

The 6.2× gap holds even when the label is missing

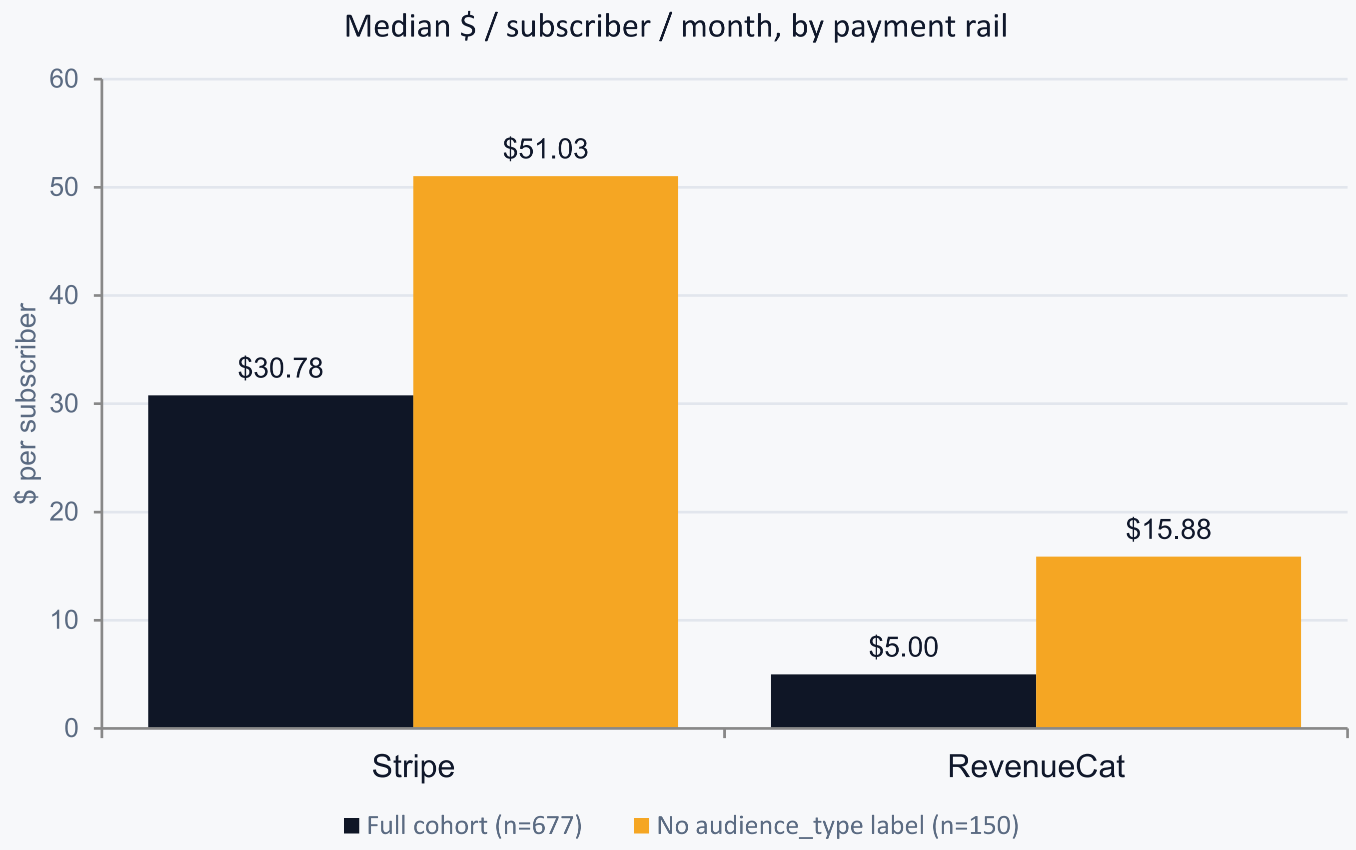

Why this works as a proxy

- RevenueCat exists to bill in-app mobile subscriptions, stating 96% of RevenueCat listings with a disclosed audience are B2C
- Stripe is general-purpose and spans both segments, but it's still where the higher-priced, lower-volume businesses sit
- Among the 183 mrr-positive listings with no audience_type at all, the same gap reappears: $51.03 (Stripe) vs $15.88 (RevenueCat), a 3.2× spread, recovered from a field that is never blank

$/subscriber computed per row as mrr_usd ÷ active_subscriptions, then medianed within group

Domain rating tracks revenue. X followers mostly don't.

Two columns claim to measure the same thing. How much of an audience a startup has built. Ranked against MRR, they diverge sharply: **domain_rating** (a proxy for accumulated SEO/backlink authority) correlates with revenue nearly three times as strongly as **x_followers** does

$\rho = 0.49$

domain_rating vs. mrr_usd (n=714)

$\rho = 0.17$

x_followers vs. mrr_usd (n=984)

Both hold up when split by segment

Domain rating's edge over followers isn't a B2B or B2C artifact. The same gap shows up inside both: B2B ρ=0.50, B2C ρ=0.39, both well clear of what followers manage anywhere

Spearman rank correlation, chosen over Pearson given the heavy right skew in both revenue and reach

2 EVIDENCE

One climbs in a straight line. The other doesn't climb at all

By domain rating quintile

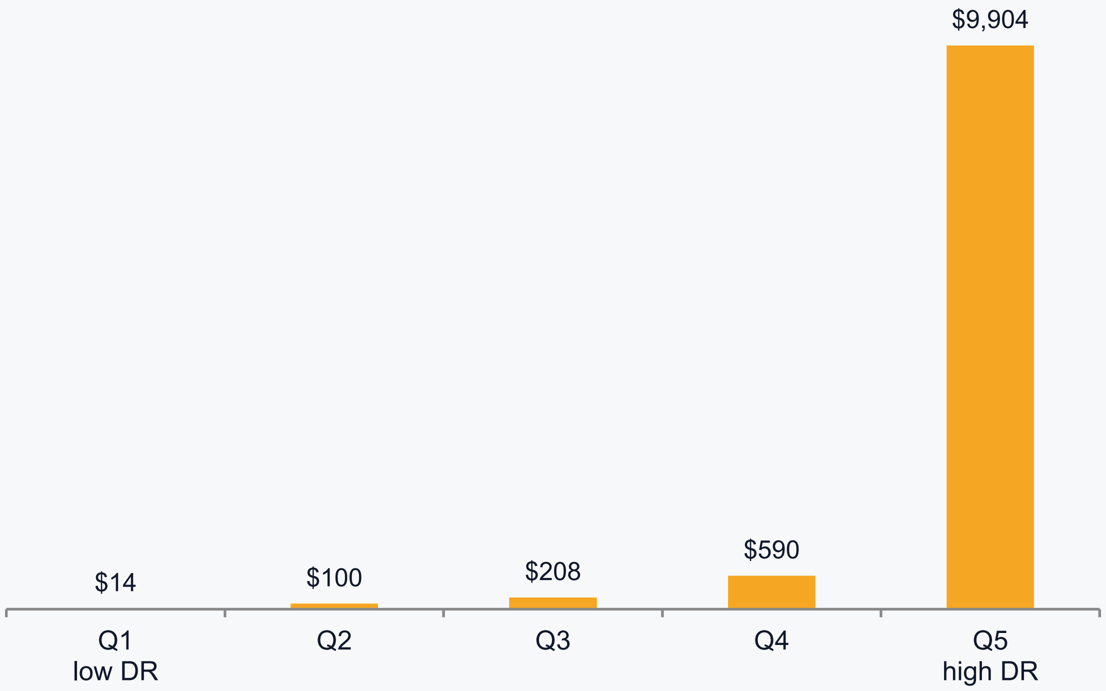

A clean staircase. Each quintile of accumulated SEO authority out-earns the last, top vs. bottom a 730× spread

By X-follower quintile

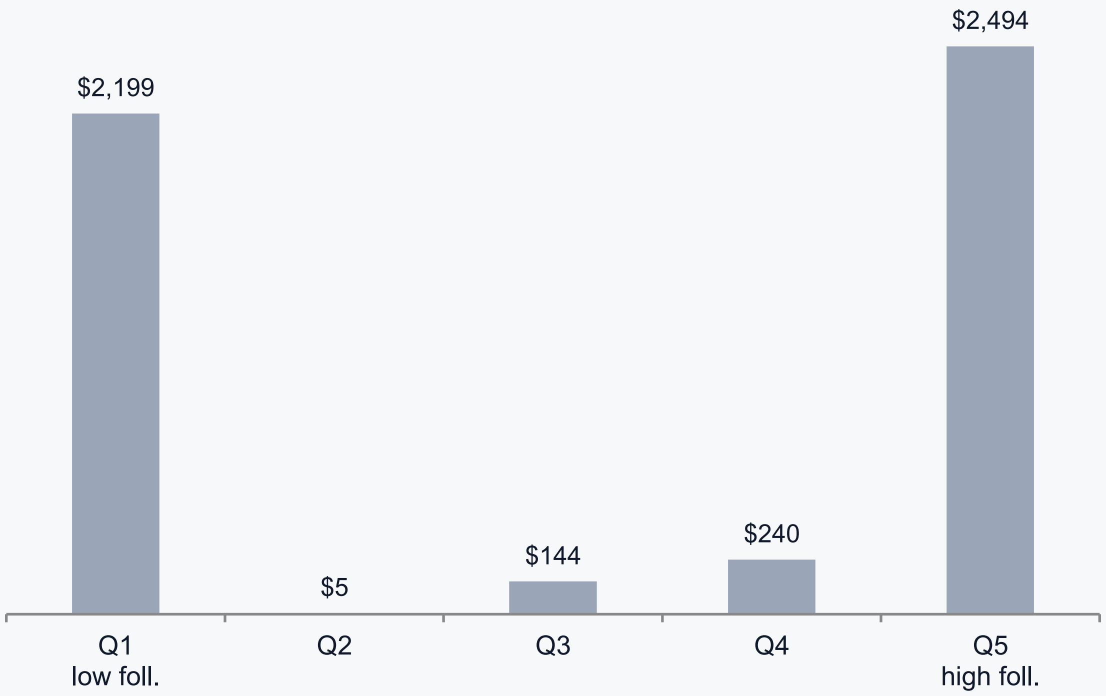

No order at all. The highest-follower quintile barely edges the lowest, and Q2 collapses to near zero

Quintiles by rank, ties broken by row order (pandas qcut). Cohorts: DR n=714, followers n=984

The listed revenue multiple looks fake until you remove the businesses too small to have one

262 listings (26.2%) are marked **on_sale**, each carrying a self-reported **revenue_multiple** alongside an **asking_price_usd** that, divided by annualised MRR, should reproduce roughly the same number. Raw correlation between the two says they don't agree at all

$r = 0.06$

stated multiple vs. implied multiple, raw (n=224)

$\rho = 0.80$

after excluding listings under $200 MRR (n=170)

The multiple isn't fabricated. It's unstable at the tail. A $20 swing in monthly revenue moves an already-tiny denominator enough to send the implied multiple from 2× to 700×

implied_multiple = asking_price_usd ÷ (mrr_usd × 12). stated multiple parsed from revenue_multiple

3 EVIDENCE

Below $200 MRR, the ratio breaks, not the marketplace

Stated vs. implied multiple agreement

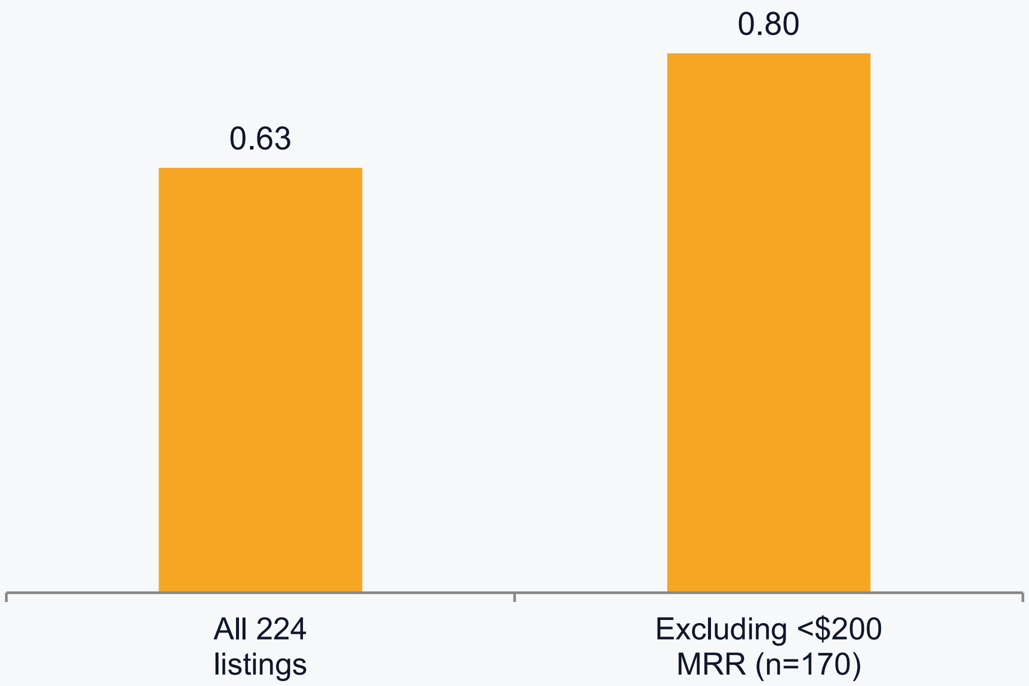

Where it breaks worst

Listing	MRR	Asking	Stated	Implied
MyBite	$1.20	$10,000	0.0x	694.4x
Rolely	$16.67	$4,000	326.8x	20.0x
Stealth Company	$54.07	$7,000	172.6x	10.8x
Abun.com	$274.00	$50,000	172.0x	15.2x

Every extreme mismatch sits under $300 in monthly revenue. The arithmetic, not the honesty, is what fails there

Practical read: ignore a listed multiple below ~$200 MRR. Recompute it from asking_price and MRR instead and above that line, the two numbers agree closely enough to trust

Extreme-mismatch examples ranked by |implied − stated| among listings with asking_price_usd > $500

SO WHAT

The self-reported fields aren't the reliable ones

Across all three signals, the field a founder fills in by hand is the weaker one. The field generated by their own infrastructure or arithmetic is the stronger one

For a buyer diligencing a listing

Don't take audience_type or revenue_multiple at face value. Cross-check payment_provider for segment, and recompute the multiple yourself once MRR clears roughly $200 while below that, treat the stated number as noise

For TrustMRR as a platform

Payment rail and computed multiple are more trustworthy than the free-text fields sitting next to them. A confidence flag on low-MRR multiples, and a rail-derived audience tag, would close two credibility gaps without asking founders for more data

For a founder reading their own metrics

A large X following is not evidence of monetizable reach and domain authority is a far better predictor of what your traffic is actually worth. Multiples quoted on near-zero MRR (yours or a comp) aren't meaningful until revenue scales past micro

How these three signals were built, and where to be careful

Derived, not disclosed fields

price_per_subscriber, implied_multiple and all quintile buckets are computed from raw columns for this analysis, but none exist in the source file

Spearman over Pearson

Revenue, followers and asking price are all heavily right-skewed; rank correlation is reported wherever a handful of outliers would otherwise dominate a linear correlation

Coverage varies by insight

Insight 1 draws on 677–984 rows depending on the cut; Insight 2 on 714 (domain_rating) and 984 (x_followers) rows; Insight 3 on the 224 on_sale listings with both a stated and computable multiple

Multiples use MRR × 12, not trailing 12-month revenue

The two diverge for volatile businesses; a revenue_12m-based multiple was checked and shows the same low-MRR breakdown

No causal claims

Payment rail, domain rating and listing price are all associated with revenue, not shown here to cause it. Selection into each rail or ranking is itself likely revenue-linked