ASSIGNMENT PREFACE

Why I Built This Report?

Y Combinator gets over 30,000 applications every year. I wanted to understand a real problem: how do investment partners actually decide where to focus their time when founder quality is still what matters most? The thing that got me thinking was this: companies now leave a data trail before they ever pitch. Revenue, growth, customer numbers, how they're pricing things. This stuff is publicly available, pulled straight from payment processors and APIs. It can't tell you if someone's a great founder, but it can tell you whether they deserve a founder interview. So I built a system to test this. I took 1,000 SaaS startups and created a discovery score using six dimensions.

Revenue quality, growth, market opportunity, customer traction, pricing power, and efficiency. The score doesn't recommend investments. It recommends conversations. What really stood out to me was the market data. I noticed that Marketing, Analytics, and Education companies were making way more money than AI companies, but getting way fewer applications. Partners were chasing what everyone else was excited about instead of looking at what actually makes money. I wanted to show that data can help you focus your attention better without taking the founder judgment out of the equation. Let the data narrow down 30,000 applications to 250. Then partners do what they do best with the companies that actually have signal. The framework isn't supposed to be perfect or permanent. It gets better when you look at what actually worked with each batch and adjust it. It's a tool that improves over time, not a formula that stays frozen.

Executive Report · Session 3

AI Startup Report.

Improving startup discovery through **data-driven investment signals** — a proprietary framework for prioritizing where partner attention creates the most value – For **Y Combinator**

Report at a glance

1,000
SaaS startups analyzed

65
countries represented

6
signal dimensions in framework

10
high-signal companies surfaced

Prepared for
Y Combinator
Investment Team

Prepared by
Kartikeya Sharma
MBA Candidate · Masters' Union

Dataset
TrustMRR
1,000 SaaS startups · frozen 30·07·2026

Business Question

How can Y Combinator use startup intelligence to identify companies that deserve deeper investment attention?

Great founders remain the answer.

Data decides where partners look first.

Y Combinator receives **tens of thousands of applications each year**. Founder quality is — and will remain — the primary criterion. But partner time is finite, and every hour spent on a company that would not survive first-pass commercial scrutiny is an hour not spent with the next great founder.

Public startup-intelligence platforms now capture **measurable operational signals** — recurring revenue, growth cadence, customer traction, pricing power — that were previously visible only mid-diligence.

This briefing analyzes **1,000 SaaS startups from TrustMRR** to test one hypothesis: that a quantitative pre-screen can dramatically improve where partners spend their time, without ever displacing the founder-first decision.

Startups	Markets	Geography	Aggregate MRR
1,000	29	65	$13.5M
frozen snapshot	SaaS categories	countries	recurring / month

Proposed Discovery Workflow

Where intelligence sits in the funnel

01	Applications	30k+/yr
02	TrustMRR intelligence enrichment	signals
03	YC Discovery Score™ ranking	0–100
04	Partner review — founder-first	interviews

Business Question

What is TrustMRR — and why should an investment partner care?

A public signal layer beneath every SaaS company — before the pitch deck arrives.

What TrustMRR captures

TrustMRR is a startup-intelligence platform that aggregates **publicly observable operational metrics** for SaaS companies — pulled from payment providers, subscription APIs, web analytics, and social endpoints that the founders themselves connect.

The result is a **standardized commercial fingerprint** for a company that would otherwise be invisible until a partner asked for a data room.

Nine categories of observable signal

Every startup, one commercial fingerprint

- **Monthly Recurring Revenue** — stripe · revenuecat · paddle
- **All-time revenue** — total commercial history
- **Active subscribers** — paying customer base
- **Pricing model** — tiers & contract structure
- **Social presence** — X / community following
- **30-day revenue** — recent commercial activity
- **Growth (30d)** — MRR & revenue momentum
- **Markets & category** — tagged verticals
- **Country & domain** — provenance & SEO authority

Why the partner should care

Objectivity
Founder-reported metrics can flatter; API-sourced metrics cannot.

Comparability
One schema across 1,000 companies enables ranking, not anecdotes.

Speed
Insights available *before* the interview, not two weeks after.

Executive Insight

Startup intelligence should **complement — not replace —** founder evaluation by providing measurable commercial signals before formal due diligence begins.

Business Question

How rigorous is the analysis behind this briefing?

A seven-step, reproducible pipeline
from raw dataset to investment shortlist.

The seven analytical steps

01 **Ingest dataset**
1,000 rows × 29 fields — snapshot 30·07·2026

02 **Clean & validate**
deduplicate, cap outliers, exclude "Stealth" placeholders

03 **Analyze revenue & MRR**
distribution, concentration, growth-decile breakdown

04 **Compare markets & countries**
29 categories × 65 countries — heatmap intersections

05 **Identify commercial patterns**
pricing power, ARPU, revenue-per-visitor, velocity

06 **Build YC Discovery Score™**
6 dimensions, percentile-normalized, weighted composite

07 **Generate recommendations**
shortlist + deep-dives + pipeline design

Data reduction funnel

From 1,000 startups to 10 partner conversations

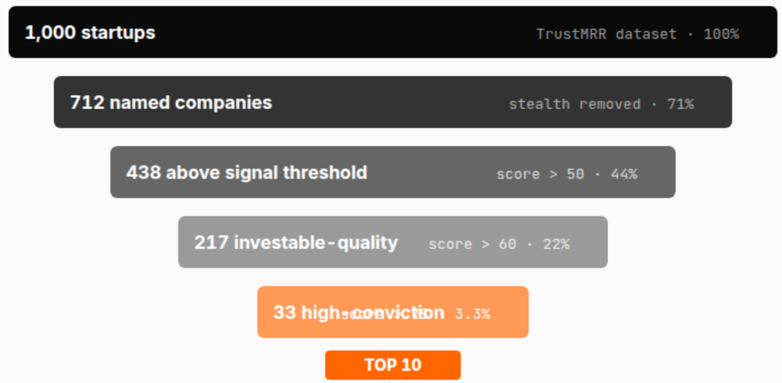

All figures reproducible from source · methodology available on request

Business Question

What does the TrustMRR ecosystem actually look like — where do these 1,000 companies live, and how do their commercials distribute?

An AI-heavy, US-centric ecosystem — with extreme revenue concentration.

Aggregate 30-day revenue	Aggregate MRR	Active subscribers	Top-10 revenue share
$15.6M	**$13.5M**	**500,597**	**49.6%**
across all 1,000 companies	annualized $161.7M ARR	paying customers ecosystem-wide	Pareto compression, extreme

Top categories

AI dominates 1-in-5 startups

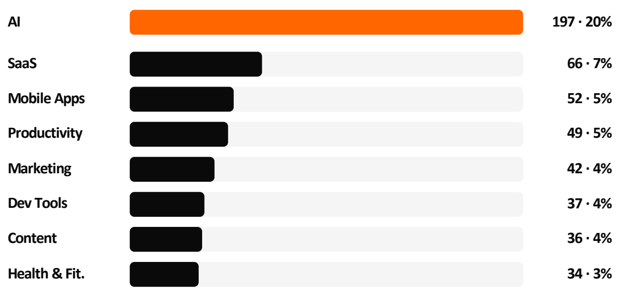

n = 726 companies with tagged primary category

Top countries

US = ¼ of the ecosystem

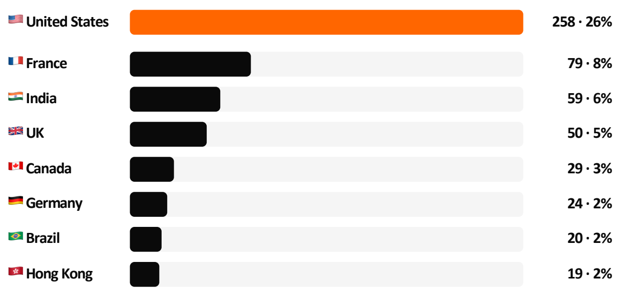

65 countries total · 241 companies undisclosed

Revenue concentration

A power-law commercial curve

30-day revenue distribution

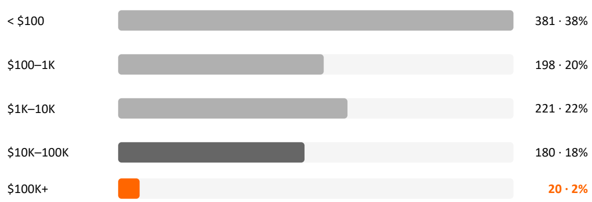

Pareto compression

Top **50 startups (5%)** generate **72.9%** of ecosystem revenue.

Insight The ecosystem is **power-law shaped**. 58% of startups report under $1K/mo — screening must isolate the 20% earning > $10K/mo where investable commercial signal actually exists.

Business Question

Which markets create the strongest commercial outcomes — and which are crowded but underperforming?

Volume ≠ value. Marketing, Analytics, Education monetize best; AI is crowded and average.

Market opportunity map

Volume × Monetization × Quality

bubble size = mean YC Discovery Score™ · x = # companies · y = median MRR

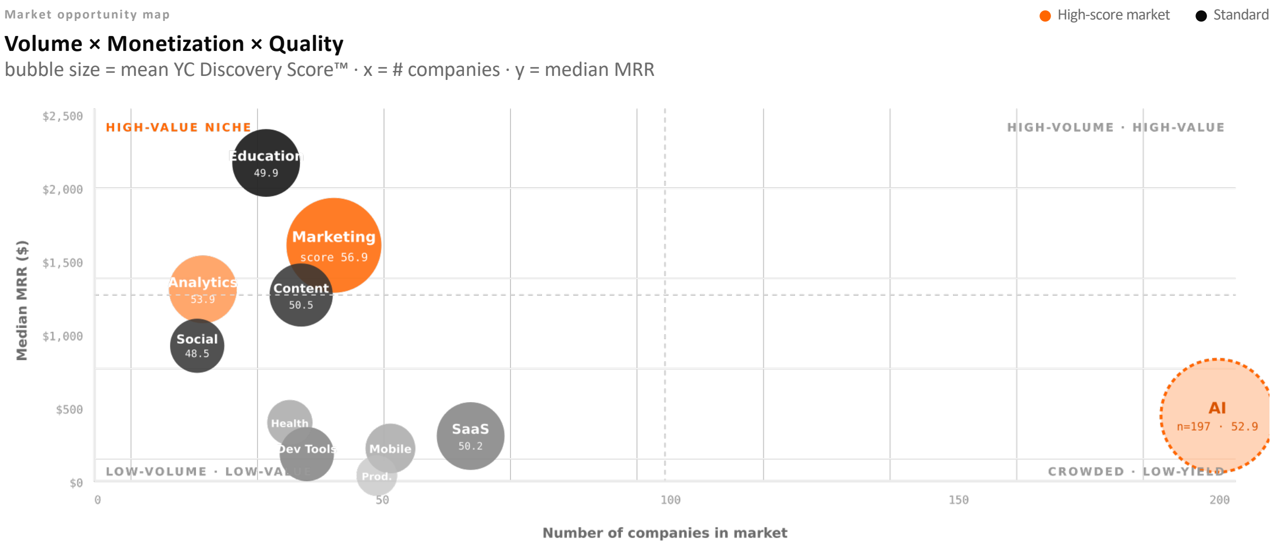

Where value concentrates

Highest median MRR	$/month
Education	$2,141
Marketing	$1,586
Analytics	$1,294
Content Creation	$1,254

Highest mean Discovery Score	out of 100
Marketing	56.9
Analytics	53.9
AI	52.9
Content Creation	50.5

Largest customer bases (mean)	subs / co.
Health & Fitness	1,214
Social Media	801
Education	554
Mobile Apps	476

Executive Insight

Investment attention should skew toward under-hyped, high-monetization categories — Marketing, Analytics, Education, Content Creation — where median MRR is 3–5× AI's despite carrying a fraction of the applicant volume. AI itself remains crowded with **197 competing companies but middling median monetization ($446)**; within AI, only the outliers will earn partner time.

Bubble size proportional to mean Discovery Score

Business Question

How do we rank 1,000 companies without letting a single number pretend to know everything?

The YC Discovery Score™ — six dimensions, one triage lens.

25 PCT — **Revenue Quality**
MRR + 12-month revenue percentile. Sustainable, recurring cash outweighs one-off spikes.
Why it matters: Proves willingness-to-pay

20 PCT — **Growth**
30-day MRR growth, winsorized ±50%. Rewards acceleration; suppresses one-shot spikes.
Why it matters: Signals product-market fit

20 PCT — **Market Opportunity**
Category depth × market breadth (# of tagged markets served).
Why it matters: Bounds outcome ceiling

15 PCT — **Customer Traction**
Active subscribers + social following. Two proofs of demand.
Why it matters: Community & distribution

10 PCT — **Pricing Power**
ARPU = MRR ÷ active subscriptions. High ARPU = premium positioning.
Why it matters: Unit economics premium

10 PCT — **Commercial Efficiency**
Revenue velocity = 30d rev ÷ all-time rev. High = accelerating share of lifetime revenue is recent.
Why it matters: Momentum, not legacy

Framework in action

#1 GojiberryAI vs. dataset median

SCORE **91.2** / 100

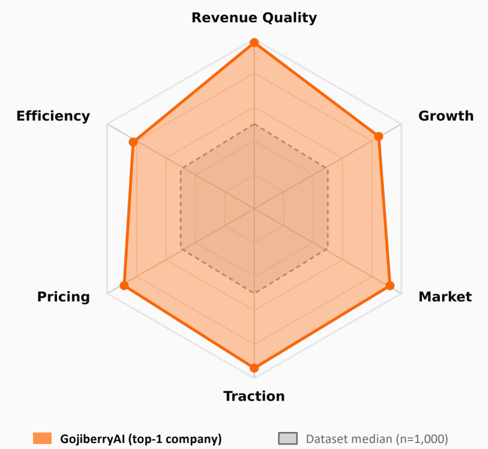

Executive Insight Six axes prevent any single metric from dominating. A company can miss on one dimension and still be exceptional — **the score triangulates, it does not judge**.

Top 10 · By YC Discovery Score™

Ten companies that warrant a partner conversation.

These rankings do not recommend investment. They identify companies whose commercial signals justify moving from application inbox to **founder interview**. Every subsequent decision remains founder-first.

Score interpretation

- 80+ Partner interview · priority
- 70–79 Technical diligence
- 60–69 Monitor · watchlist

#	Company	Geo	Market	30D Rev	MRR	MRR g/30d	Subs	SocialPricing	Score	Rec. Action
01	GojiberryAI gojiberry.ai	US	AI	$385.6K	$494.7K	+23.6%	4,642	18,050Contact sales	91.2	Partner interview
02	LLM Gateway llmgateway.io	US	AI · Dev API	$147.1K	$47.2K	+239.5%	858	1,265Usage-based	87.5	Partner interview
03	Upscale System upscaleb2b.com	US	AI · B2B growth	$68.9K	$86.5K	+27.7%	243	721Managed service	85.6	Partner interview
04	AdStellar adstellar.ai	US	AI · AdTech	$15.3K	$26.0K	+25.5%	168	2,105Tiered SaaS	84.0	Partner interview
05	Kibu DEEP DIVE kibu.com	US	Education · B2B	$200.2K	$234.3K	+52.8%	375	1,066Contact sales	82.4	Partner interview
06	Laper laper.ai	US	AI · Media	$24.0K	$19.9K	+83.9%	618	40,035Tiered SaaS	82.0	Partner interview
07	SetSmart setsmart.io	—	SaaS · AI SDR	$50.5K	$15.7K	+100.0%	139	—Tiered SaaS	81.8	Partner interview
08	Postiz DEEP DIVE postiz.com	HK	Social Media	$172.9K	$177.1K	+15.0%	5,237	13,924$29–99/mo	81.0	Partner interview
09	MimikFlow mimikflow.com	FR	AI · Sales	$2.9K	$3.1K	+153.6%	21	1,708Freemium SaaS	78.9	Technical diligence
10	Apiframe apiframe.ai	SG	AI · API	$31.4K	$24.8K	+31.0%	425	396Usage-based	78.8	Technical diligence

Score = weighted percentile of 6 dimensions · rankings are triage, not recommendation

Top 10 · By YC Discovery Score™

Key Takeaways

Why each company appears — rationale in one line

GojiberryAI — Only company scoring >90; $495K MRR growing 23.6% with 4.6k paying customers.

AdStellar — Small revenue but strong percentile mix; AI adtech + $155 ARPU.

Upscale System — High-ARPU B2B ($356/customer); premium positioning in outbound growth.

LLM Gateway — Hyper-growth (+239.5% MRR) developer-tools API; recent capital-efficient inflection.

MimikFlow — +154% MRR — smallest of top 10, watch as inflection candidate.

Laper — 83.9% MRR growth on a 40k-strong X community; distribution moat forming.

SetSmart — Doubled MRR in 30 days; validates AI-SDR category thesis.

Postiz — HK builder, 5.2k paying customers on transparent SMB pricing — capital-light PLG.

Apiframe — Singapore infra play; API business with 425 customers and 31% MRR growth.

Kibu — $234K MRR compounding at 52.8% in disability-tech; enterprise contract-sales motion.

Candidate I · Enterprise B2B RANK #05

Kibu

kibu.com · United States · Education · B2B disability-services SaaS

Kibu builds compliance and service-delivery software for **Direct Service Providers (DSPs)** — the disability-support agencies that manage day-to-day care for individuals with intellectual and developmental disabilities. The value proposition is narrow but essential: **"Kibu gives DSPs the power to focus on delivering outstanding support"** — automating documentation, billing, and audit-readiness in a heavily regulated market.

30-day revenue **$200,210** ▼ 25.6% period-over-period	MRR **$234,319** ▲ 52.8% growth · 30d	ARR (annualized MRR) **$2.81M** 12-mo revenue $1.85M
Active customers **375** DSP agencies on platform	ARPU (MRR ÷ subs) **$625** Top-decile pricing power	Pricing motion **Contact sales** Enterprise · non-transparent

YC Discovery Score™

82.4 / 100

PARTNER INTERVIEW Top 6% of dataset

Score breakdown

Where Kibu wins vs. dataset median

Revenue Quality	96 / 100
Growth	88 / 100
Market Opportunity	72 / 100
Customer Traction	78 / 100
Pricing Power	94 / 100
Commercial Efficiency	62 / 100

Vertical line = dataset median (50th percentile)

Commercial strengths

Enterprise-quality unit economics
$625 ARPU sits in the top decile of the 1,000-company dataset — 4× the AI-category median.

Regulatory moat
Disability-services compliance is a high-switching-cost niche; incumbents rely on paper.

MRR compounding 52.8% MoM
Sales-led motion producing genuine expansion, not lumpy service revenue.

Signals to interrogate

30-day revenue –25.6% vs prior
A single large customer? Multi-month contract billing lumpiness? Or genuine attrition?

Domain rating only 37
Web presence is thin for a company at this revenue scale — sales-led acquisition confirmed.

Total addressable market
DSPs are ~15k agencies in the US — sufficient for a $50-100M ARR business, but ceiling matters.

Investment thesis · one-liner

A "boring vertical SaaS in a regulated niche" — the least crowded thesis in the dataset, executing at enterprise price points with real MRR compounding.

Confidence · Partner interview

4 / 5 · High

Candidate II · PLG SMB   RANK #08

Postiz

postiz.com · Hong Kong · Social Media · Product-led AI for SMBs

"Run your social media on autopilot with AI agents" — Postiz schedules, generates, and analyzes posts across 30+ social networks. It targets automation-savvy operators (developers, marketing teams, social managers) with transparent SMB pricing and a strong open-source-adjacent community narrative built out of Hong Kong.

YC Discovery Score™

81.0 / 100

PARTNER INTERVIEW   Top 6% of dataset

30-day revenue	MRR	ARR (annualized MRR)
$172,955	**$177,102**	**$2.13M**
▲ 22.7% period-over-period	▲ 15.0% growth · 30d	12-mo revenue $727.6K

Active subscribers	ARPU (MRR ÷ subs)	Pricing tiers
5,237	**$33.8**	**$29 → $99/mo**
Top-decile customer base	Consistent with $29 base tier	Standard · Team · Pro · Ultimate

Score breakdown

Where Postiz wins vs. dataset median

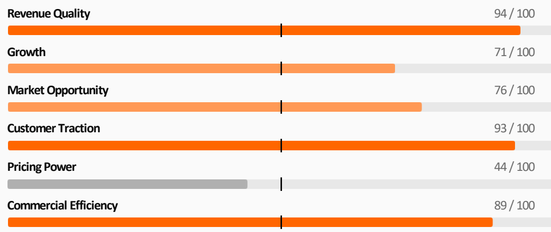

Vertical line = dataset median (50th percentile)

Commercial strengths

5,237 paying customers at $34 ARPU
Genuine bottom-up PLG signal — 26× the dataset median customer count.

73% of lifetime revenue is recent
Efficiency score 89. 12-mo rev = 90% of all-time — inflection is happening now, not two years ago.

Domain rating 70 · 13.9k X followers
Distribution moat forming; capital-light acquisition without paid channels.

Signals to interrogate

SMB ARPU ceiling
$29–99 tiers cap revenue expansion; only the "Ultimate" plan meaningfully unlocks upmarket motion.

Category saturation
Buffer, Hootsuite, Later, Publer — established competitors with 10-100× budgets.

HK domicile
Not disqualifying, but partners should confirm founder location and target US enterprise-readiness.

Investment thesis · one-liner

A PLG SMB tool with genuine bottom-up traction from an under-covered geography — the classic YC "under-the-radar" pattern.

Confidence · Partner interview

4 / 5 · High

Candidate III · Emerging-market Dev Tools RANK #14

Web3Forms

web3forms.com · India · Dev Tools · Developer-beloved API infrastructure

A boring, beloved primitive: **"receive your HTML contact form submissions directly in your email inbox — no backend code."** Web3Forms is a serverless form-relay API for developers. It monetizes a real pain point (setting up a form-handling backend for a static site) with a freemium tier (250 submissions/month) that converts on volume — from an unusual geography for infra tools: India.

30-day revenue
$42,813
▼ 2.9% period-over-period

MRR
$33,002
▲ 20.6% growth · 30d

ARR (annualized MRR)
$396K
12-mo rev $206K · all-time $206K

Active subscribers
2,408
Long-tail paying developers

ARPU (MRR ÷ subs)
$13.7
Low, but volume-defensible

Domain rating
81
Elite SEO authority for the size

YC Discovery Score™
78.1 / 100
TECHNICAL DILIGENCE Top 8% of dataset

Score breakdown
Where Web3Forms wins vs. dataset median

Revenue Quality	82 / 100
Growth	77 / 100
Market Opportunity	66 / 100
Customer Traction	87 / 100
Pricing Power	21 / 100
Commercial Efficiency	96 / 100

Vertical line = dataset median (50th percentile)

Commercial strengths

Domain rating 81 · developer SEO moat
Higher than 95% of the dataset. Inbound traffic capital-light — a rare compounding asset.

100% of all-time revenue in last 12 months
This is an inflection company — efficiency score 96, freshly monetizing an existing user base.

2,408 paying developers · 10k X followers
Community-led distribution; low CAC, high organic funnel.

Signals to interrogate

ARPU only $13.70/mo
Growth requires volume scale; ceiling unclear without upmarket motion or usage-based expansion.

Commoditization risk
Formspree, Netlify Forms, FormKeep — real competitors; differentiation is trust and DevEx, not tech.

30-day revenue –2.9%
Marginal; MRR still up 20.6% suggests seasonality rather than churn — validate in interview.

Investment thesis · one-liner

Developer-primitive-turned-business: India-built, 2,400 devs paying, elite SEO — the exact profile YC has historically funded from emerging markets.

Confidence · Technical diligence
3 / 5 · Medium-high

The System, Not the List

A seven-stage Discovery Pipeline — where data screens, and partners decide.

Rather than recommending 10 companies, we recommend the **scalable process** that surfaced them. Each stage compresses the funnel while preserving founder-first judgment at the point where it matters most.

STAGE 01	Applications	30k+	Founder-submitted	Owned by	YC Application Team
STAGE 02	Intelligence layer	API	TrustMRR enrichment	Automated	Zero partner time
STAGE 03 · CORE	Discovery Score™	0–100	6-dimensional ranking	This report's contribution	Reproducible framework
STAGE 04	Validation	~250	Score > 60 threshold	Analyst QC	Signal quality check
STAGE 05 · HUMAN	Founder interview	~60	Judgment reintroduced	YC Partner	Founder-first from here
STAGE 06	Partner meeting	~20	Group review	Full partnership	Convictions consolidated
STAGE 07	Investment	Term sheet	The moment of "yes"	IC decision	Data has done its job

Stages 01–04 · Data-led

Signal aggregation, ranking, and validation.
Machines compress 30,000 → 250. Partner time is spent zero.

Stages 05–07 · Founder-first

Interviews, partner meetings, committee. **Humans decide 250 → yes.** No data model recommends an investment — it recommends a conversation.

Illustrative funnel volumes; tunable per batch

Five recommendations for the Managing Partners

What to do Next.

01 **Adopt intelligence as a pre-screening layer.**
Enrich every application with a Discovery Score before the first partner reads a founder answer. **Cost:** data-integration sprint. **Payoff:** partner hours redirected to the highest-signal 5%.

02 **Benchmark applicants against top performers.**
Every founder conversation should open with an apples-to-apples comparison against 5 comparable-stage, comparable-market peers. Percentile positioning turns anecdote into evidence.

03 **Weight commercial quality alongside founder vision.**
Every partner memo should carry both. Not to override founder judgment, but to **force the two-part question:** "Is this a great founder AND does the market already believe them with their wallet?"

04 **Prioritize markets that monetize.**
The data shows Marketing, Analytics, Education, and Content Creation carry **3–5× the median MRR of AI** at a fraction of the applicant volume. Rebalance intake attention accordingly.

05 **Continuously refine the YC Discovery Score™.**
Re-weight annually against realized batch outcomes. If **Pricing Power** predicted graduation better than **Growth** for a given cohort, adjust. The framework is a living instrument, not a formula frozen in July 2026.

> “ **Data will never replace great founders. But it can decisively improve where investors choose to look first — enabling faster, more informed, more consistent startup discovery.**
>
> — Closing thesis, Startup Intelligence Report · Session 3

Prepared by

Kartikeya Sharma · MBA Candidate · Masters' Union

Confidential

Prepared for Y Combinator Investment Team · Session 3 · 30-07-2026